

Plasticity in Deep Reinforcement Learning

Submitted to the
Albert-Ludwigs-Universität Freiburg

by Mohammad Mahdi Gheidi



Master's Project

Submitted on **August 11, 2026**
to the **Albert-Ludwigs-Universität Freiburg**
by **Mohammad Mahdi Gheidi** (5973722)
Examined by **Prof. Dr. Joschka Bödecker**

Neurobotics Laboratory
Department of Computer Science (IIF)
Faculty of Engineering
Albert-Ludwigs-Universität Freiburg

Acknowledgements

I would like to thank my supervisor Dr. Shengchao Yan for his support throughout this project and letting me pursue what I find meaningful with this work.

I also thank my family for their love and support throughout the way.

Abstract

Deep reinforcement learning has achieved remarkable success in domains such as game playing, robotics, recommendation systems, control, and large-scale decision making [Mnih et al., 2015, Silver et al., 2017]. However, despite these successes, modern deep reinforcement learning systems frequently suffer from a phenomenon known as plasticity loss.

Plasticity in general is the ability of a system to change how it perceives and processes the data and be able to adapt. Considering the human brain and the studies performed on it, neuroplasticity is the brain's remarkable capacity to reorganize, grow, and rewire its neural connections throughout an individual's lifetime. Rather than being a fixed organ, it continuously adapts to new experiences, learning processes, and environmental stimuli, or following an injury [Puderbaugh and Emmady, 2023]. In the context of deep learning and neural networks, it refers to a neural network's ability to adapt its predictions when new information becomes available. As training progresses, networks often become increasingly difficult to modify, resulting in slower learning, reduced adaptability, and performance degradation in non-stationary environments [Lyle et al., 2023, Ash and Adams, 2020].

In recent years, we have seen an increasing rate in the adoption of RL systems in various domains, from self-driving cars to autonomous robots and superhuman results on many games. As we use DRL systems more and more by the day, we need to find the root causes preventing neural networks from benefitting of plasticity while solving tasks in the reinforcement learning regime.

The objective of this project is to conduct a literature review and investigate the mechanisms responsible for plasticity loss and understand how this phenomenon affects deep reinforcement learning systems. We base our benchmark on the reproduction of the results from [Lyle et al., 2023] and argue about the effect of geometrical landscape of functions and their effect on learning.

Contents

1	Motivation	1
2	Fundamentals	2
2.1	Stability-Plasticity Dilemma	2
2.2	Primacy Bias	2
2.3	Plasticity Loss	3
3	Methodology	4
3.1	Neural Network Architectures	4
3.1.1	Multi-Layer Perceptron (MLP)	4
3.1.2	Convolutional Neural Network (CNN)	4
3.1.3	Vision Transformer (ViT)	4
3.2	Datasets	4
3.3	Experimental Environments	5
3.3.1	EasyMDP	6
3.3.2	HardMDP	6
3.3.3	Sparse-reward MDP	6
3.4	Measuring Plasticity	6
3.5	Loss Landscape Analysis	8
3.6	Deep Q-Network Learning	10
4	Results	12
4.1	Optimizer Instability Under Abrupt Label Changes	12
4.2	Hessian Spectrum and Gradient Covariance	13
4.3	Falsification of Simple Explanations	15
4.4	Learning Curve Evolution on Probe Tasks	15
4.5	Effect of Network Width on Plasticity Loss	16
4.6	Effect of Plasticity-Preserving Interventions	16
5	Conclusion	19
	Bibliography	I

1 Motivation

Plasticity is becoming increasingly important as reinforcement learning systems are deployed in environments that continuously evolve.

Unlike traditional supervised learning, reinforcement learning operates under fundamentally non-stationary targets, especially when using value-based learning methods such as learning Q-values.

As the learning progresses during an RL task, policies and state values continuously change, data distributions evolve during training and the targets moves as learning objectives depend on previous updates e.g. temporal-difference learning methods. This creates a non-stationary learning problem, in which we are chasing a moving target.

As tasks become more complex, the ability of neural networks to remain adaptable becomes increasingly important. For example on long-horizon tasks and continuous-action environments there is no other feasible solution than to use neural networks. In that case we need to make sure that this tool is a reliable function approximator to learn the dynamics of a changing environment.

Several recent works suggest that plasticity loss may be one of the major factors limiting the scalability of deep reinforcement learning. This has been noticed before in the deep learning regime, e.g. in the works purely done on supervised learning methods, we see that training the model on some part of the dataset first and training it again on the other parts of the same dataset results in poor generalization [Ash and Adams, 2020]. Other works in DRL have shown that, for example, as off-policy methods rely on using a replay buffer and use the data to train a model, will create the issue that the model overfits to early experiences in a way that harms future learning [Nikishin et al., 2022].

Understanding and preventing this phenomenon would lead to:

- More stable RL algorithms
- Faster adaptation
- Better continual learning
- Improved sample efficiency
- More robust autonomous agents

2 Fundamentals

The roots of this work can be traced to several related research areas.

2.1 Stability-Plasticity Dilemma

A long-standing challenge in deep reinforcement learning - and more generally in neural networks of all kinds - is balancing stability and adaptability. Systems must preserve useful knowledge while remaining capable of learning new information.

A trade-off we often encounter in our implementations of the neural networks is the balance between stability and plasticity. Having control and deeper understanding of this phenomenon is crucial for developing robust RL systems. As we often would like to keep the previous knowledge and at the same time learn new information, we need to find a balance between stability and plasticity.

We know that having excessive stability leads to rigidity and inability to learn on new signals. A network which shows frozen-like behavior and is unable to adapt to new information is not useful in a non-stationary environment. On the other hand, having excessive plasticity leads to catastrophic forgetting.

This issue would not be much of a concern in stationary environments, in which the data distribution does not change over time. However, in non-stationary environments, the ability to adapt to new information is crucial for success.

2.2 Primacy Bias

Neural networks have been developed based on the idea of a neural system of the human brain. Primacy bias, also known as the primacy effect was introduced as the idea of cognitive bias that describes the tendency for humans to remember information presented at the beginning of a sequence better than information presented later on. This bias has been observed in various contexts, including memory recall, decision-making, and learning. For example on learning a new language, if we happen to learn the pronunciation of a word incorrectly at first, we will have a hard time to correct it later on. This is because the initial information we absorb has a stronger influence on our memory and perception than subsequent information. That is why we often have a hard time to correct our mistakes and learn new information. This bias can have significant implications for learning and decision-making, as it can lead to the formation of inaccurate or incomplete mental representations of information.

Neural networks show the same tendency to be influenced by early experiences, which can lead to a bias in the learned representations and affect the network's ability to learn from new information. [Nikishin et al., 2022] have already demonstrated that early experiences exert disproportionately large influence on neural network representations.

Networks become increasingly biased toward information learned early in training, reducing their ability to adapt later.

2.3 Plasticity Loss

Building on what we have already discussed, we can define plasticity loss as the loss of the ability of a neural network to adapt to new information over time. This phenomenon is particularly relevant in the context of deep reinforcement learning, where agents must continuously learn and adapt to changing environments.

[Ash and Adams, 2020] showed that networks can become increasingly difficult to retrain after prolonged optimization. Their observations provide evidence that optimization itself changes trainability. In the following sections we show that a trained network will have a harder time to adapt to new learning signals, compared to a network that has been freshly initialized. This is a clear indication of the primacy bias in neural networks, which can and will effectively lead to plasticity loss and reduced adaptability over time.

The most comprehensive recent study presented as a benchmark, which also discusses some remediation methods is [Lyle et al., 2023].

This paper provides strong evidence that plasticity loss is fundamentally related to changes in the geometry of the loss landscape rather than simple explanations, e.g. dead neurons, low capacity or other previously mentioned improvements on the matter.

In the next section we cover the methodology used to carry out these experiments. We've based this project mainly on the ideas adapted from this paper.

3 Methodology

In order to be able to understand the underlying mechanisms causing the observed phenomena we need to define a methodology to observe under-the-hoods of the training procedure of neural networks.

3.1 Neural Network Architectures

To evaluate the effects of different neural network designs on plasticity, three commonly used architectures are considered throughout this project: a Multi-Layer Perceptron (MLP), a Convolutional Neural Network (CNN), and a Vision Transformer (ViT). These architectures were selected because they represent three fundamentally different approaches to feature extraction and representation learning, allowing the investigation of whether plasticity-related phenomena are architecture-dependent.

3.1.1 Multi-Layer Perceptron (MLP)

The MLPs we consider in our implementation consists of two fully connected hidden layers. Unless otherwise specified, each hidden layer contains 512 neurons. For the network width scaling experiment presented later in this report, the hidden layer width is varied to investigate the relationship between model capacity and plasticity. In these experiments, a base width of 16 neurons is multiplied by scaling factors of 1, 2, 4, 8, 12, and 16, resulting in progressively wider networks.

3.1.2 Convolutional Neural Network (CNN)

The CNN architecture consists of two convolutional layers followed by two fully connected layers. The first convolutional layer employs 5×5 kernels, while the second uses 3×3 kernels. Both convolutional layers contain 64 feature maps. The extracted features are then processed by two fully connected layers, each containing 256 hidden units before producing the final output.

3.1.3 Vision Transformer (ViT)

The Vision Transformer implementation follows the architecture proposed by [Dosovitskiy et al., 2020]. Images are first divided into patches of size 3×3 , which are embedded into a feature representation before being processed by a transformer encoder. The model uses an embedding dimension of 256, a feed-forward network dimension of 1024, a single transformer encoder block, and a dropout rate of 0.1. Unlike models initialized from pretrained weights, all parameters are trained from scratch during every experiment.

3.2 Datasets

We used the MNIST and CIFAR-10 datasets throughout the experiments. These datasets are among the most widely used benchmarks for image classification and provide two complementary levels of visual complexity. The MNIST dataset consists of 70,000 grayscale images




Figure 1: CNN on MNIST.

of handwritten digits (60,000 training and 10,000 testing images), where each image has a resolution of 28×28 pixels and belongs to one of ten digit classes (0–9).

The CIFAR-10 dataset contains 60,000 colour images (50,000 training and 10,000 testing images), each with a resolution of 32×32 pixels and belonging to one of ten object categories, including airplanes, automobiles, birds, cats, deer, dogs, frogs, horses, ships, and trucks. Compared to MNIST, CIFAR-10 presents a substantially more challenging classification task because of its greater visual diversity, colour information, and more complex object features. Using both datasets enables the evaluation of neural network plasticity under learning tasks of varying complexity while maintaining a consistent ten-class classification setting.

Table 1: Datasets used throughout the experiments.

Dataset	Image Size	Channels	Classes	Training / Test
MNIST	28×28	1 (Grayscale)	10	60,000 / 10,000
CIFAR-10	32×32	3 (RGB)	10	50,000 / 10,000

3.3 Experimental Environments

To carry out the experiments in this project, we do not use a large-scale reinforcement learning benchmark, but rather toy environments to create controlled datasets in which plasticity can be measured reliably.

To achieve this, the methodology proposed by [Lyle et al., 2023] converts standard supervised learning datasets into Markov Decision Processes (MDPs). This design provides a controlled setting where the effects of plasticity loss can be isolated from other challenges commonly encountered in reinforcement learning.

The experiments use the MNIST and CIFAR-10 datasets and consider three environment variants.

3.3.1 EasyMDP

EasyMDP corresponds to the true-label setting.

Each image is associated with its original dataset label. The objective of the agent is therefore identical to standard classification. The environment serves as a baseline in which the network learns meaningful representations from naturally structured targets.

3.3.2 HardMDP

HardMDP corresponds to the random-label setting.

Before training begins, a random label is assigned independently to each image. Consequently, images belonging to the same class may receive entirely different labels.

For example, one image of the digit 7 may be assigned label 5, while another image of the 7 may be assigned label 2.

This environment removes the underlying structure of the dataset and forces the network to memorize arbitrary associations. The setting is intentionally difficult and is used to investigate how plasticity evolves under extreme memorization demands.

3.3.3 Sparse-reward MDP

Sparse-reward MDP corresponds to the reinforcement learning setting used in the original paper.

In this environment, the agent receives a reward only when it selects the correct action in a designated goal state. In the experiments, state 9 is treated as the rewarding state.

The agent must therefore learn value estimates for all states and propagate reward information through the value function. Unlike EasyMDP and HardMDP, this environment introduces sparse rewards and temporal-difference learning, thereby capturing the essential characteristics of reinforcement learning.

3.4 Measuring Plasticity

As mentioned above, the primary objective of this work is not to maximize reinforcement learning performance or model accuracy, but rather to understand and quantify the phenomenon of plasticity loss. Therefore, a precise operational definition of plasticity is required.

Following [Lyle et al., 2023], plasticity is defined as the ability of a neural network to update its predictions in response to new learning signals. More formally, consider an optimization algorithm:

$$\mathcal{O} : (\theta, \ell) \mapsto \theta^*, \quad (1)$$

which takes a parameter vector θ and a loss function ℓ and returns updated parameters θ^* . The resulting parameters need not correspond to a global optimum; for example, $\mathcal{O}$ may simply perform a fixed number of gradient descent steps.

To measure the flexibility of a network, a distribution over learning objectives is considered. In particular, [Lyle et al., 2023] define a family of regression objectives

$$\ell_{f,X}(\theta) = \mathbb{E}_{x \sim X} \left[(f(\theta, x) - g_\omega(x))^2 \right], \quad (2)$$

where g_ω defines a randomly generated target function.

To construct challenging prediction targets that require the network to modify its predictions in arbitrary directions, the target function is transformed according to

$$g(x) = a + \sin(10^5 f(x; \omega_0)), \quad (3)$$

where ω_0 is sampled from the same initialization distribution as the network parameters.

The offset parameter a is chosen to match the mean prediction of the network. This prevents the probe task from favouring randomly initialized networks whose outputs tend to be centered around zero.

The following definition is used to quantify this adaptability [Lyle et al., 2023]. Let $\mathcal{L}$ denote a distribution over probe loss functions. For a network checkpoint θ_t and a sampled probe objective $\ell \sim \mathcal{L}$, the optimization procedure produces the adapted parameters

$$\theta_t^* = \mathcal{O}(\theta_t, \ell), \quad (4)$$

where $\mathcal{O}$ is the same optimization procedure used for the network’s primary learning objective.

The adapted parameters θ_t^* need not be a global optimum; they are simply the parameters obtained after a fixed optimization budget.

Let b denote the loss achieved by a baseline predictor. For a regression objective, b may, for example, be chosen as the variance of the target values. The plasticity of the network at checkpoint θ_t is then defined as

$$\mathcal{P}(\theta_t) = b - \mathbb{E}_{\ell \sim \mathcal{L}} [\ell(\theta_t^*)], \quad \theta_t^* = \mathcal{O}(\theta_t, \ell). \quad (5)$$

A network is considered more plastic when it can obtain a lower loss on the sampled probe objectives within the fixed optimization budget. Consequently, a larger value of $\mathcal{P}(\theta_t)$ indicates greater plasticity.

Following Lyle et al. [2023], the change in plasticity relative to initialization is defined as

$$\Delta \mathcal{P}_t = \mathcal{P}(\theta_t) - \mathcal{P}(\theta_0), \quad (6)$$

where θ_0 denotes the randomly initialized network. Because the baseline b appears in both terms, it cancels:

$$\Delta \mathcal{P}_t = \mathbb{E}_{\ell \sim \mathcal{L}} [\ell(\theta_0^*) - \ell(\theta_t^*)]. \quad (7)$$

Under this signed convention, $\Delta \mathcal{P}_t = 0$ indicates that plasticity has not changed relative to initialization, $\Delta \mathcal{P}_t < 0$ indicates plasticity loss, and $\Delta \mathcal{P}_t > 0$ indicates improved adaptability. More negative values therefore correspond to more severe plasticity loss.

The resulting objective measures how easily a trained network can adapt its predictions to new targets. Networks capable of rapidly minimizing this objective are considered highly plastic, whereas networks that struggle to adapt are considered to have lost plasticity.

3.5 Loss Landscape Analysis

To better understand the optimization behaviour and plasticity of neural networks, this work analyzes the geometry of the loss landscape throughout training. In particular, two complementary quantities are investigated: the Hessian matrix of the loss function and the gradient covariance matrix. Together, these metrics provide insight into the curvature of the optimization landscape

and the similarity of gradient directions across different training samples.

The Hessian matrix describes the second-order curvature of the loss function with respect to the network parameters. For a neural network with parameters θ and loss function $\ell(\theta)$, the Hessian is defined as

$$H_\ell(\theta) = \nabla_\theta^2 \ell(\theta) \in \mathbb{R}^{d \times d}, \quad (8)$$

where $d = |\theta|$ denotes the total number of trainable parameters. Since the Hessian is generally a very large matrix, its eigenspectrum is analyzed instead,

$$\Lambda(H_\ell(\theta)) = (\lambda_1, \lambda_2, \dots, \lambda_d), \quad \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d. \quad (9)$$

The eigenvalues quantify the curvature of the loss landscape along different directions in parameter space, while the corresponding eigenvectors indicate these principal directions of curvature. The largest eigenvalue is commonly used as a measure of the sharpness of the loss landscape, whereas the ratio between the largest and smallest eigenvalues (the condition number) provides an indication of the numerical conditioning of the optimization problem. Poor conditioning generally results in slower convergence and more difficult optimization.

In addition to local curvature, this work also studies the similarity between gradients computed from different training samples. Let $x_1, \dots, x_k$ denote a randomly sampled mini-batch of training examples. The gradient covariance matrix is computed as

$$C_k[i, j] = \frac{\langle \nabla_\theta \ell(\theta, x_i), \nabla_\theta \ell(\theta, x_j) \rangle}{\|\nabla_\theta \ell(\theta, x_i)\| \|\nabla_\theta \ell(\theta, x_j)\|}, \quad (10)$$

which corresponds to the cosine similarity between the gradients of every pair of training samples. Consequently, every entry of C_k lies within the interval $[-1, 1]$.

Positive values indicate that two samples produce similar parameter updates, suggesting cooperative learning and shared feature representations. Conversely, negative values imply that the corresponding gradients point in opposing directions, meaning that improving the performance on one sample may degrade the performance on the other. Such conflicting updates are commonly referred to as *gradient interference* and have been identified as an important factor affecting plasticity and continual learning.




Figure 2: Illustration of a loss landscape together with the principal eigenvectors of the Hessian matrix. The eigenvectors indicate the directions of maximum and minimum curvature, while the corresponding eigenvalues quantify the curvature along those directions. Image generated using the generative model GPT5.6.

Finally, the rank and overall structure of the gradient covariance matrix are examined throughout training. A low-rank covariance matrix indicates that many gradients lie within a small subspace of the parameter space, revealing shared optimization directions across the dataset. Visualizing this matrix as a heatmap further exposes clusters of highly correlated or anti-correlated gradients, providing additional insight into how the network organizes its internal representations as learning progresses.

3.6 Deep Q-Network Learning

A Deep Q-Network (DQN) [Mnih et al., 2015] is used throughout the experiments.

Although exact dynamic programming methods could be applied to represent the Q-values directly, such approaches would not allow investigation of the internal dynamics of neural network learning. Since the objective of this project is to study plasticity loss within neural networks, a function approximator is required.

The DQN learns action values according to the Bellman equation

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a'). \quad (11)$$

The network parameters are optimized using stochastic gradient descent on the temporal-difference loss

$$L(\theta) = (y - Q(s, a; \theta))^2, \quad (12)$$

where

$$y = r + \gamma \max_{a'} Q(s', a'; \theta^-) \quad (13)$$

and θ^- denotes the target network parameters.

4 Results

Using the tools available at our disposal, we have done 6 experiments. The first 6 experiments of the original paper are reproduced and above we have summarized the results and our insights from these experiments. There are 3 main questions we wish to gain more intuition on. First, we would like to see what would be the result of the networks losing plasticity. In an essence, what is the reason we care and what are the repercussions of this phenomenon. In the first experiment we show that the optimizer instability under abrupt label changes will lead to a large fraction of inactive ReLU units and thus loss of representational capacity. Then we move to the second experiment and show that the Hessian spectrum and gradient covariance are affected by the optimization process and thus the geometry of the loss landscape is changed in a way that makes it harder to optimize on new objectives.

Second question we wish to answer, is to what properties would cause plasticity loss? We try to investigate the effect of long-discussed weight norm and dead-units and show that there's no direct causal relationship between these metrics and loss of plasticity. Then we move to our other experiment and try to measure the plasticity of trained networks after arbitrary learning time and see that learning speed reduces as training proceeds.

At the end we have put our efforts to verify the solutions to alleviate this issue, by trying to answer several questions. Such as would increasing the size of the network be enough, or which kind of interventions in the networks is the most beneficial. We see that unlike other regimes, in which increasing the capacity will help, in deep reinforcement learning world that we have non-stationary targets, simply increasing the size of the network has not that effect. Then we analyze the effect of few previously suggested methods, such as shrink and perturb, resetting last layer, layer-normalization and others on two architectures of simple MLPs and Vision Transformers; We have observed that the paper's claim is indeed relevant and the most effective remedies are layer normalization and categorical representations (or also known as two-hot encoding).

As a disclaimer, we would like to point out that the compute resources available to us are limited and we have to make some compromises on the size of the networks and the number of experiments we can run. We will try to reproduce the results from [Lyle et al., 2023] as closely as possible. On some of the more compute-extensive tasks like the ones involving CNNs and Vision Transformers, we will have to reduce the number of experiments and seeds. Thus we may see an increase in variance in our results compared to the original study.

4.1 Optimizer Instability Under Abrupt Label Changes

The purpose of this experiment is to study how an adaptive optimizer behaves when the learning objective changes abruptly. In this setting, the network is first trained on MNIST images with randomly assigned labels. After the network has learned this mapping, the labels are randomly reassigned while the input images remain fixed.

This creates a sudden change in the underlying gradient structure of the network. The optimizer must now adapt to a new objective that is inconsistent with the previously learned one. In the case of Adam, the first- and second-moment estimates no longer accurately describe the new gradient distribution. As a result, the effective update sizes can become unstable. The experiment also shows a tuned version of Adam, which tries to mitigate this sudden change by adding a small constant.

The reproduced experiment shows that this instability can cause a large fraction of hidden ReLU units to become inactive. Once these units die, the network loses effective representational capacity and its ability to learn the new task is strongly reduced. This provides an example of plasticity loss caused by optimizer dynamics under non-stationarity.




Figure 3: Abrupt label changes destabilize Adam and lead to a large fraction of inactive ReLU units.

4.2 Hessian Spectrum and Gradient Covariance

The goal is to compare how the loss landscape evolves under gradient-based optimization versus random perturbations of the parameters. We use hessian curvature and gradient covariance to analyze the loss landscape. In both cases, the magnitude of the parameter change is controlled, but the direction of the update differs.

The Hessian spectrum is used to analyze the curvature of the probe loss landscape. Larger spectral outliers indicate sharper directions in parameter space and therefore potentially worse optimization conditioning. The reproduced Hessian analysis shows that gradient-based training tends to produce larger Hessian outliers than random Brownian perturbations. This suggests that training does not merely move the network away from initialization, but moves it in directions that make the local loss landscape harder to optimize.

Gradient covariance matrices are computed to analyze interference between samples. Since




Figure 4: Hessian spectra and gradient covariance matrices, in the beginning and after 20 epochs of training.

the entries of the gradient covariance matrix are cosine similarities between normalized gradient vectors, their values must lie in the interval $[-1, 1]$. Therefore, in the reproduced plots we use the corrected color range $[-1, 1]$ rather than the wider range shown in the original visualization $[-5, 5]$. In the gradient covariance heatmaps, positive values indicate aligned gradients, while negative values indicate gradient interference. We observe that after training for 20 epochs, the gradient covariance matrix shows more negative interference between samples compared to the initial state, meaning that the network cannot change its representation independently on many samples and directions. This suggests that gradient-based optimization can create a less favorable structure for adapting to new objectives, which may contribute to plasticity loss.

For clarity, we note that the Hessian and gradient covariance are computed using the probe loss, which is a separate loss function used to evaluate the network’s ability to adapt to new objectives. This allows us to isolate the effects of training on the loss landscape from the specific task being learned.

4.3 Falsification of Simple Explanations

This experiment investigates whether previously proposed quantities such as weight norm and number of inactive units can provide a causal explanation of plasticity loss.

The key idea is that a true causal mechanism should show a consistent relationship with plasticity across different environments. However, the reproduced results show that these quantities do not behave consistently. In some settings, the correlation with plasticity loss is positive, while in others it is negative or weak. Therefore, these quantities may be correlated with plasticity loss in particular cases, but they do not provide a general causal explanation.




Figure 5: Previously proposed explanations do not show a consistent relationship with plasticity loss across various environments.

4.4 Learning Curve Evolution on Probe Tasks

The purpose of this experiment is to understand whether plasticity loss corresponds to poor local minima or to slower optimization progress on new objectives.

The reproduced learning curves show that networks from later training checkpoints reduce the probe loss more slowly than networks near initialization. This suggests that plasticity loss is expressed primarily as slower adaptation, rather than simply convergence to a worse final plateau.




Figure 6: Later checkpoints show slower optimization on probe tasks, indicating reduced plasticity.

4.5 Effect of Network Width on Plasticity Loss

This experiment studies whether increasing network width is sufficient to prevent plasticity loss.

The reproduced results indicate that wider networks generally suffer less plasticity loss than smaller networks. However, increasing width alone does not fully eliminate the problem. This suggests that plasticity loss is not merely a consequence of insufficient model capacity, but is also related to optimization dynamics and the geometry of the loss landscape.

4.6 Effect of Plasticity-Preserving Interventions

This experiment compares several interventions designed to reduce plasticity loss, including layer normalization, reset-based methods, shrink-and-perturb, weight decay, spectral normalization, and categorical output representations.

The reproduced results support the conclusion that interventions which improve the smoothness and conditioning of the loss landscape are the most effective. In particular, layer normalization and categorical output representations show strong improvements compared to methods that only perturb or regularize the parameters. This supports the broader hypothesis that plasticity loss is closely related to loss landscape geometry.




Figure 7: Increasing width reduces plasticity loss but does not fully remove it.




Figure 8: Interventions that smooth or stabilize the loss landscape are most effective at preserving plasticity.

5 Conclusion

This project investigated the mechanisms behind plasticity loss in neural networks, with a particular focus on deep reinforcement learning. The main objective was to understand why networks that have already undergone training often become less capable of adapting to new learning signals. To address this question, we studied the recent literature on plasticity loss, capacity loss, warm-starting, and value-function learning, and reproduced several central experiments from [Lyle et al., 2023] to verify and get the understanding of what the authors have conveyed in this seminal work of theirs.

After reproducing these experiments in this project, we can support the view that plasticity loss is closely connected to the geometry of the loss landscape. In particular, the optimizer instability experiment showed that abrupt changes in the learning objective can significantly disturb the gradient structure and cause many hidden ReLU units to become inactive. This prevents the network from effectively using its representational capacity and can strongly reduce future learning ability. The Hessian and gradient covariance experiments further showed that training can move the network into regions of parameter space with sharper curvature and stronger gradient interference. These observations suggest that plasticity loss is not only a representational problem, but also an optimization problem.

Based on the reproduced experiments and the methods studied in [Lyle et al., 2023], the most effective interventions were those that made loss curvature easier to optimize. In particular, methods such as layer normalization and categorical output representations were more successful than methods that only perturbed or regularized the parameters. This supports the hypothesis that preserving plasticity requires controlling the sharpness and conditioning of the loss landscape.

Layer normalization is especially relevant in this context because it changes the parameterization of the network in a way that can stabilize optimization. By normalizing intermediate activations, layer normalization can reduce sensitivity to scale changes and help prevent the network from entering extremely sharp regions of the loss landscape. In the experiments of [Lyle et al., 2023], the addition of layer normalization reduced plasticity loss in several architectures. This suggests that normalization methods may help preserve plasticity by producing smoother and better-conditioned optimization dynamics.

The results are also consistent with the conclusions of [Farebrother et al., 2024], who argue that value functions in deep reinforcement learning should not necessarily be trained as standard regression problems. In conventional value-based RL, the network is trained to regress directly to bootstrapped target values. This can be difficult because the targets are noisy, non-stationary, and unbounded. Regression objectives may therefore push the parameters into regions of the loss landscape from which later adaptation becomes difficult. In contrast, categorical methods such as two-hot encoding or HL-Gauss transform value learning into a classification-style

problem. Instead of directly predicting a scalar value, the network predicts a distribution over a fixed support. This can make the learning problem better conditioned and more stable.

Further works can study methods to use to reduce the sharpness or improve the network's navigation through function space that plasticity remains intact.

Bibliography

Jordan Ash and Ryan P Adams. On warm-starting neural network training. *Advances in neural information processing systems*, 33:3884–3894, 2020.

Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020.

Jesse Farebrother, Jordi Orbay, Quan Vuong, Adrien Ali Taïga, Yevgen Chebotar, Ted Xiao, Alex Irpan, Sergey Levine, Pablo Samuel Castro, Aleksandra Faust, et al. Stop regressing: Training value functions via classification for scalable deep rl. *arXiv preprint arXiv:2403.03950*, 2024.

Clare Lyle, Zeyu Zheng, Evgenii Nikishin, Bernardo Avila Pires, Razvan Pascanu, and Will Dabney. Understanding plasticity in neural networks. In *International Conference on Machine Learning*, pages 23190–23211. PMLR, 2023.

Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex Graves, Martin Riedmiller, Andreas K Fidjeland, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015.

Evgenii Nikishin, Max Schwarzer, Pierluca D’Oro, Pierre-Luc Bacon, and Aaron Courville. The primacy bias in deep reinforcement learning. In *International conference on machine learning*, pages 16828–16847. PMLR, 2022.

Matt Puderbaugh and Prabhu D Emmady. Neuroplasticity. In *StatPearls [Internet]*. StatPearls Publishing, 2023.

David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, et al. Mastering the game of go without human knowledge. *Nature*, 550(7676):354–359, 2017.